Rethinking AI Disclosure Across Content Type and Creator–Reader Relationships

Hyewon Lee*
School of Computing, KAIST
Daejeon, Republic of Korea
hyewon0809@kaist.ac.kr

Ihchae Ryu*
School of Computing, KAIST
Daejeon, Republic of Korea
ihchaeryu@kaist.ac.kr

Yumin Cho*
School of Computing, KAIST
Daejeon, Republic of Korea
dbals1184@kaist.ac.kr

Hwajung Hong
Industrial Design, KAIST
Daejeon, Republic of Korea
hwajung@kaist.ac.kr

Abstract

As AI-generated content becomes increasingly prevalent, discussions around the necessity and design of AI disclosure have gained critical importance. While prior studies have primarily focused on user perceptions of AI disclosure or explored disclosure design within specific content types, little is known about how disclosure expectations differ across content genres and between creators and readers. To address this gap, we conducted a formative survey (N=33) investigating disclosure perceptions across seven text content types, revealing noticeable mismatches between creator and reader expectations depending on the context. Building on these insights, we conducted in-depth group interviews (N=6) using participatory activities to explore how disclosure preferences differ and how they are negotiated through interaction. Our findings demonstrate that AI disclosure perceptions are shaped not only by content characteristics, but also by the relational dynamics between creators and readers. Moreover, disclosure expectations were found to impact behaviors throughout the content lifecycle, from creation to evaluation. This work highlights the need for flexible, context-sensitive, and relationship-aware approaches to AI disclosure design.

CCS Concepts

• **Social and professional topics** → **Computing / technology policy**; • **Human-centered computing** → **Empirical studies in HCI**; **User studies**.

Keywords

AI disclosure, Content type, Creator–Reader relationship, User perception, Transparency

ACM Reference Format:
Hyewon Lee, Ihchae Ryu, Yumin Cho, and Hwajung Hong. 2025. Rethinking AI Disclosure Across Content Type and Creator–Reader Relationships. In *Unpublished manuscript*. ACM, New York, NY, USA, 9 pages.

1 Introduction

In recent years, artificial intelligence—especially in the form of generative AI—has been transforming how content is produced and consumed. Tools like Gemini[1] and DALL · E[2] have made it easy to generate high-quality content across modalities such as text, image, and videos, for both personal and professional purposes. As the scope and sophistication of generative AI expands, the media landscape is rapidly shifting: distinguishing between human-created and AI-generated content is becoming increasingly difficult, if not impossible [3]. This growing ambiguity has raised serious concerns about transparency, misinformation, and deception, especially since people often assume content is human-made by default [11]. In response, there is a global movement toward mandating the disclosure of AI use in media—an effort reflected in legislative initiatives such as the EU AI Act[15] and the US AI Disclosure Act[16].

Despite this momentum, there remains a lack of consensus on how AI disclosure should be designed and communicated. While previous work has called attention to the importance of transparency, the specific methods, formats, and contextual considerations involved in disclosure remain underexplored. In relevance, a recent workshop [4] identified critical questions that need to be addressed in themes such as policy, trust, and user experience. Moreover, the notion of 'AI-generated' is ambiguous: its interpretation varies depending on the degree and type of AI involvement, the perception of the user and also the terminology and design of the disclosure [5, 9]. To be meaningful and actionable, disclosure design must account for these layers of complexity that shape both the production and perception of content.

This paper investigates how people perceive AI disclosure in relation to different types of text-based content, focusing especially on the dynamic between content creators and readers. Through a formative survey study (N=33), we explore how disclosure expectations and preferences differ across content types, and between those who write and those who read them. Our results reveal clear mismatches in the expectations of creators and readers about the disclosure of certain types of content. To further unpack these discrepancies, we conducted in-depth interviews and participatory activities with paired creator-reader participants (N=6) across selected content scenarios. This multi-phase study surfaces the contextual factors that contribute to perceptual gaps and highlights aspects of AI disclosure that may be negotiable versus inherently contested. We argue that effective AI disclosure must be content-sensitive and attuned to creator-reader relationships and interactions. Our

*Both authors contributed equally to this research.

Preprint, KAIST, Daejeon, Republic of Korea
2025.

[1] https://gemini.google.com
[2] https://openai.com/index/dall-e-3/

study contributes actionable insights that can guide effective disclosure design and foster further discussions essential for shaping a future where AI is deeply interwoven with human culture and communication.

2 Related Work

2.1 Effects of AI Disclosure toward Assessment on Content and Creator

Prior studies consistently demonstrate that individuals tend to devalue content once they encountered AI disclosure. For example, in a study comparing argumentative essays and creative stories, participants rated writing quality significantly lower when AI assistance was disclosed [8]. Similarly, in the domain of prosocial advertising, disclosing AI involvement led to less favorable reader perceptions, which subsequently influenced donation behaviors [1]. These effects are also pronounced in emotionally evocative content, such as poetry, where AI disclosure provokes negative evaluations [2, 13].

More specifically, AI disclosure has been found to influence key dimensions of content evaluation, including perceived trustworthiness, credibility, and accuracy [1, 7, 10]. Beyond the content itself, AI disclosure also shapes how people perceive the human creator: notably, creators are often attributed with less effort and diminished achievement when their work is known to involve AI [12].

In this study, we seek to explore the underlying reasons behind the assessment of AI-disclosed content, examining both the perspectives of readers and creators. By doing so, we aim to expand the discussion on AI disclosure, making it more practical and relevant for real-world contexts.

2.2 AI Disclosure Design and Its Effects

Beyond the binary question of whether AI involvement should be disclosed, recent studies have explored how the format, timing, and granularity of disclosure shape user perception and how such disclosures should be designed.

Lim et al. [9] emphasize that AI disclosure is not merely a matter of transparency but a critical factor that affects content evaluation. Their study further shows that a reader's preexisting perception toward AI interacts with disclosure, jointly influencing how content is perceived. Several studies have also examined how different disclosure label designs impact users' perceptions of content and their subsequent actions. Gamage et al. [6] tested various visual and textual label designs for AI-generated images on social media, while Epstein et al. [5] studied user responses to AI-generated news headlines. These works highlight how subtle design elements—such as tone, placement, and wording—can significantly affect credibility judgments and user engagement. However, these studies primarily focus on the reader's perspective, paying less attention to the relational dynamics between creators and readers or how disclosure is interpreted from the creator's point of view.

Resnik and Hosseini [14] address AI disclosure strategy in the context of scientific writing. They argue that AI involvement should be disclosed only when it is intentional and substantial, excluding trivial use cases such as typo correction or citation search. Their framework offers a nuanced approach to classifying AI usage in the creating process but still falls short in capturing diverse content contexts.

Our work extends this conversation by emphasizing that before converging on specific disclosure designs, it is essential to first explore and facilitate mutual understanding and negotiation between creators and readers across diverse content contexts shaped by AI involvement.

3 Formative Study

We conducted a formative study using an online survey to demonstrate how perceptions of AI disclosure vary across different content categories and depend on the roles of both creators and readers.

3.1 Participant

The survey was distributed through an online community and conducted in Korean, open to anyone with the provided link. Over three days, 33 participants completed the survey. Among them, 18 individuals reported frequent AI use in reading or writing text content, rating 4 or 5 on a 5-point Likert scale. Most participants were in their 20s (M = 27.12, SD = 9.41), with four participants in their 40s and 50s. The survey took approximately 21 minutes to complete. Five participants were randomly selected to receive a $3.50 gift card as compensation.

3.2 Procedure

Content Type	Description
News	Article-style content that delivers the latest updates and information.
Email	Messages exchanged online between computer users.
Social Media Post	Public posts shared online, such as Instagram posts or personal blogs.
Novel	Prose based on imagination or real events written by an author.
Personal Statement	A document introducing oneself, including name, background, and career.
Textbook	Instructional material used in schools for academic learning.
Assignment (Essay)	A structured piece of writing on a specific topic, often logical and analytical.

Table 1: Seven content types selected for the formative study.

The survey consisted of six questions related to seven diverse text contents (Table 1), incorporating both multiple-choice and short-answer formats. The authors selected these seven contents to ensure a broad range of content categories (e.g., informational, personal, creative, professional, educational) that individuals commonly encounter in their daily lives. All contents included the same set of questions (refer to the Appendix for details).

3.3 Results

3.3.1 Preferences for AI Disclosure by Content Type. We asked participants what information should be included in AI disclosures when content is created with AI. Multiple selections were allowed. Responses revealed two distinct preference patterns. For *Email*, *Social Media Post*, and *Personal Statement*, participants predominantly favored simply indicating whether AI was used, without further detail (54.5%, 45.5%, and 45.5%, respectively; Table 3). In contrast, for *News*, *Novel*, *Textbook*, and *Assignment (Essay)*, participants showed a strong preference for specifying where and when AI was used (69.7%, 69.7%, 78.8%, and 78.8%, respectively; Table 4).

3.3.2 Obligations and Preferences for AI Disclosure from Creator and Reader Perspectives. Participants were asked to evaluate AI disclosure obligations and preferences from both the creator and reader perspectives across different content types. Based on the degree of alignment between the two perspectives, the results were categorized into three groups: strong match, medium match, and low match (Figure **??**).

In the strong match category, which included *Email* and *Textbook*, the dominant opinions of creators and readers were aligned. For *Email*, both parties agreed that disclosure was unnecessary and generally undesirable. For *Textbook*, both agreed on the necessity and desirability of disclosure.

The medium match category included *News* and *Assignment (Essay)*. In these cases, both creators and readers agreed that disclosure was obligatory. However, their preferences diverged, while readers preferred disclosure to be present, creators were evenly split between wanting and not wanting to disclose.

In the low match category—*Social Media Post*, *Personal Statement*, and *Novel*—there were significant discrepancies between the two perspectives. Readers largely supported mandatory disclosure and expressed a strong preference for AI usage to be visible. In contrast, creators showed no clear consensus, with opinions on both obligation and preference being evenly divided.

3.3.3 Factors Influencing AI Disclosure Across Content Types. We also collected additional opinions on AI disclosure across different content types. The results revealed both shared and divergent influencing factors depending on the content's nature (Table 2). Notably, for *Novel*—the only creative writing content—responses were predominantly negative and aggressive. For example, P9 stated, *"I want to avoid garbage text written by AI unfaithfully,"* while P20 expressed strong disapproval toward those who use AI to profit from writing without being genuine authors. In contrast, *Textbook* was perceived differently, as they are typically produced through collaborative authorship and careful attribution of sources. P17 noted, *"Since textbooks are usually compiled from the contributions of multiple individuals, adding an AI-generated opinion would not be particularly problematic."*

4 User Study

Through the formative study, we revealed that contextual characteristics of each content type, as well as the relationship between creators and readers, can influence expectations and perceived impact of AI disclosure. Based on these findings, we formulated the following research questions:

- **RQ1**. What factors contribute to the differing perspectives between creators and readers regarding AI disclosure across different text content types?
- **RQ2**. Among these differences, which can be reconciled through negotiation, and which represent more fundamental, irreconcilable tensions?

4.1 Participants

To address the research questions, we selected three content types and conducted group interviews with one creator and one reader per content. Six participants were recruited from the formative study, all of whom had indicated willingness to participate in follow-up interviews. We selected *Personal Statement*, *Novel*, and *Assignment (Essay)* based on the degree of perspective divergence observed in the formative study. Among these, *Personal Statement* and *Novel* showed low matches, while *Assignment (Essay)* exhibited a medium match between creators and readers.

4.2 Procedure

Each study session lasted 90 minutes and focused on one of the three selected content types. Each session consisted of three phases: **Introduction** (15 minutes), **Individual Activity** (30 minutes), and **Group Discussion** (45 minutes).

During the **Individual Activity**, participants reviewed example content and its AI-assisted creation process, both prepared by the authors (see Appendix for details). They reflected on which aspects they would disclose as creators or expect to see disclosed as readers, and completed a brief design task proposing disclosure methods.

In the **Group Discussion**, participants discussed their differing perspectives as creators and readers, referencing the formative study findings and their individual activities. The session aimed to facilitate mutual understanding and negotiation of disclosure norms. A semi-structured interview was also conducted to gather deeper insights.

All sessions were conducted via Google Meet[3], with collaborative activities facilitated using FigJam[4]. Sessions were recorded with participants' consent, and each participant received a gift card equivalent to $7.25 in local currency.

4.3 Results

In this section, we present and summarize the results of each session, organized around three key themes: (1) aligned perspectives between creators and readers, (2) conflicting perspectives between creators and readers, and (3) notable discussion points.

4.3.1 Personal Statement. Creators and readers agreed that disclosure should be made when AI influences the content or authenticity of the statement. They both mentioned that authenticity is the key element in personal statements, as it is the writing that shows oneself. Minor assistance, such as length adjustment, was viewed as non-essential to disclose. However, they diverged on when and what to disclose. Creators preferred disclosing after reading to avoid bias, while readers wanted it before reading to facilitate efficient evaluation. U1 questioned why the use of AI was being judged

[3]https://meet.google.com/landing
[4]https://www.figma.com/figjam/

Content	Keyword		
	Impact Type	**Core Value**	**Other Attributes**
News	-	Credibility, Fairness	Aversion to Mass Production
Email	Personal	Sincerity, Affection, Competence	-
Textbook	Social Impact	Factuality	Co-authored
Novel	-	Sincerity, Creation, Art	Profitability
Personal Statement	-	Sincerity, Competence	Evaluation
Social Media Post	Wide-scale Influence	-	Aversion to Mass Production
Assignment (Essay)	-	Fairness	Evaluation

Table 2: Factors influencing AI disclosure across content types (keyword analysis)

more harshly than traditional forms of human assistance, such as tutoring or editing services. U2 suggested that the use of AI could itself demonstrate a person's ability or resourcefulness, and thus the process of using AI should be considered part of the evaluation of a personal statement. They also noted that discussions about evaluation criteria—specifically, what readers aim to assess through a personal statement—should precede debates about AI disclosure.

4.3.2 Novel. There was consensus between creator and reader that full disclosure is necessary if an AI-generated sentence or paragraph is used without modification. Conversely, minor AI support such as grammar checking or content feedback was considered acceptable without disclosure if the creative contribution remained with the author. However, creators and readers differed on the scope and format of disclosure. Creators wished to clarify their personal contributions, where and how AI was involved to avoid any misunderstanding, while readers focused solely on whether AI was used or not—regardless of the stage. Readers also expected a quantifiable measure of AI involvement, whereas creators found it difficult to assign numeric values to nuanced contributions. Creators were concerned with ownership and dignity in AI use, while readers prioritized immersive reading experience. They also revealed a tradeoff between transparency and narrative immersion.

4.3.3 Assignment (Essay). Both creator and reader agreed that educational systems should prioritize teaching the responsible use of AI tools rather than just prohibiting AI use. Creators believed proper citation sufficed, making AI disclosure unnecessary. In contrast, readers viewed disclosure as a matter of academic ethics and preferred it at the end of the essay to avoid bias during reading. U5 noted that, since readers are in a position to evaluate the creator's writing, a fundamental consensus on AI disclosure would be difficult to achieve unless readers' perceptions change.

5 Discussion

5.1 Content Characteristics Shape Perceptions toward AI Disclosure

Our studies revealed that perceptions toward AI disclosure vary significantly depending on content characteristics.

For *News* and *Textbook*, participants frequently mentioned terms like *public interest* and *social impact*, suggesting that AI disclosure is necessary. P32 stated, "News is public writing with authority, so the contribution of AI should be disclosed". Similarly, P8 emphasized, "In textbooks, the social impact of misinformation is significant, so it's essential to clarify how and where AI was used". In contrast, email, viewed as *private communication*, was seen as less in need of AI disclosure. P28 explained, "I don't think AI disclosure is necessary because knowing AI helped write one might feel like there's a lack of effort or personal connection."

The relationship between the creator, the reader, and the work also influenced perceptions of AI disclosure. In *Novel*, where the work itself is evaluated, U3 considered how reader evaluates the text when suggesting timing of AI disclosure. In *Personal Statement*, where the reader evaluates the creator through the work, P28 expressed concerns about disclosing AI use, fearing it might make them appear incompetent to the reader. In *Assignment (Essay)*, where the interaction between creator and work is key, U6 argued that the focus should be on whether creators are using AI appropriately, rather than on the disclosure itself.

5.2 Context within the Same Content Type Shape Perceptions toward AI Disclosure

Even within a single type of content, the context and purpose can influence how AI disclosure is perceived.

For example, the perceived needs of AI disclosure in a *Personal Statement* varied depending on contextual factors such as the evaluation criteria, position applied for, and authority level of the institution. U2 pointed out, "Readers might also want to see how the creator used AI in their writing, considering it as part of their abilities."

Similarly, in a workshop on *Novel*, both genre and medium—whether paper or digital—influenced participants' reactions to encountering AI disclosure. U4 expressed concern about AI's accuracy in historical fiction, while U3 viewed AI as an innovative tool in science fiction. Additionally, U4 suggested that "an option to toggle AI disclosure in an e-book would be helpful," but noted that such markers in paper books could be "potentially distracting."

Therefore, AI disclosure standards cannot be determined solely by content type. The perception of AI disclosure varies according to context and purpose, requiring a nuanced, context-specific approach in its design.

5.3 AI Disclosure Serves as a Space for Mutual Understanding and Negotiation

Our user study revealed that the majority of participants—four out of six (U2, U3, U4, and U6)—changed their perspectives on AI disclosure during the user study. For example, U2, a *Personal Statement* reader, reflected, "Initially, I questioned the need for disclosure. But after the study, I realized it could be interesting to know how the creator used AI."

During group discussions, participants often empathized with each other's perspectives. U4, a *Novel* reader, shared, "I thought the creator's perspective was completely different from mine, but after hearing their explanation, I realized I could relate." Similarly, U3, a *Novel* creator, noted, "Even if I were a reader, I would want to know the exact percentage of AI assistance in the writing."

Participants were sometimes surprised to find that the perspectives of other differed from their expectations. U3 remarked, "What surprised me was realizing that readers do not necessarily want as much information as I had expected."

These findings suggest that AI disclosure is not just about transparency. Rather, the act of disclosing—and the conversations it triggers—serve as a relational and ethical negotiation space, where stakeholders collaboratively interpret and assign meaning to AI involvement.

5.4 AI Disclosure Shapes Behavior Across the Content Lifecycle

Our study reveal that AI disclosure impacts both creators' and readers' behavior throughout the entire content lifecycle, extending well beyond its intended purpose of ensuring transparency in final outcomes.

Several participants reported that disclosure could affect their decision to engage with a text. U4, a *Novel* reader, commented, "If there's AI disclosure, people who are strongly against AI might refuse to read it. While for others, it could rather spark curiosity." Similarly, P14 noted, "I need AI disclosure before purchasing, because I don't want to consume a novel involving heavy AI assistance."

The disclosure methods also shape creators' behaviors before and during the creation process. U2 shared, "If I were asked to submit the entire AI chat log for writing my personal statement, I probably wouldn't apply at all." This suggests that certain disclosure requirements might deter content creation entirely. U2 also added, *"If I knew I'd have to disclose, I might run two separate AI chats—one to actually draft my personal statement, and another to generate a more presentable version that I could disclose."*

These findings suggest that AI disclosure is not a one-way mechanism that affects only content assessment. Instead, it acts as a bidirectional force that directly influences stakeholder behaviors at every stage from creation to consumption.

5.5 AI Disclosure is Tightly Linked to Social Perceptions

Consistent with previous studies, our study shows that underlying perceptions toward AI significantly influence how readers interpret AI-disclosed content. P9 noted, "AI-written news feels like it may contain inaccurate information, and I probably wouldn't read it carefully." Similarly, P15 shared, "When reading a text fully written by AI, I often feel like it's a waste of time." P20 added, "If I sense AI usage from awkward phrasing, my trust in the writer drops, as it suggests they lack editing skills."

Creators often tailor their stance on AI disclosure accordingly. U1 said, "If AI disclosure appears upfront, readers may approach the text with preconceived notions." U5 shared, "If a text is disclosed as AI-written, there's bound to be bias. I wouldn't want my work to be undervalued because of it, so I wouldn't disclose it."

However, our study also reveals that AI usage does not always lead to negative perceptions. U3 noted, "In genres like science fiction, AI involvement could even be seen as an advantage." Similarly, U2 stated, "Using AI for a personal statement can be seen as a personal skill, and I would be curious to see how much AI someone can leverage."

These findings suggest that AI disclosure cannot be defined by a single fixed policy. Even if a definitive standard that meets current requirements were to be established, such a policy would need to remain flexible and evolve alongside changing social perceptions and specific contexts.

6 Limitations and Future Work

6.1 Limitations

First, there are potential biases in participant backgrounds. The majority of participants in both studies were women in their twenties with above-average familiarity with AI. Second, participants often possessed inadequate proficiency as a creator or reader. For example, the definition of the "novel content creator" role in the survey included not only professional writers but also amateur novelists. Lastly, the AI-assisted content creation scenarios used in the user study lacked diversity, as only one scenario was considered for each content type.

6.2 Future Work

This paper focused on exploring contextual differences among text-based content. By expanding the scope to include different modalities and more specific types of content, future research can identify crucial factors directly related to content-specific AI literacy. As discussed earlier, a practical framework that enables creators and readers to understand each other and find common ground can also be developed. In addition, future work can propose concrete and effective disclosure labels or layouts through a relationship-focused approach to AI disclosure design.

7 Conclusion

Through the formative study, we demonstrate that preferences and needs regarding AI disclosure vary depending on the type, attributes, and context of the content. Through the user study, we highlight the importance of approaching AI disclosure as a relational process, one that involves negotiating a shared understanding that takes into account the multiple aspects of the creator-reader relationship.

In conclusion, for AI disclosure to be practical and effective, it must be grounded in a comprehensive understanding of content context, type, and the dynamics between creators and readers prior

to designing specific labels or systems. Also, AI disclosure should not be treated as a fixed standard, but rather as a flexible social mechanism that evolves through user engagement and negotiation. Discrepancies in expectations and perceptions can—and should—be addressed through open dialogue among stakeholders. Finally, AI disclosure significantly influences the perceptions and actions of both creators and readers. Ongoing research and observation are essential to fully understand its evolving impact.

Acknowledgments

All three authors contributed equally to this work. The order of authors is alphabetical.

References

[1] Tae Hyun Baek, Jungkeun Kim, and Jeong Hyun Kim and. 2024. Effect of disclosing AI-generated content on prosocial advertising evaluation. *International Journal of Advertising* 0, 0 (2024), 1–22. doi:10.1080/02650487.2024.2401319 arXiv:https://doi.org/10.1080/02650487.2024.2401319

[2] Hao Chen, Pingping Wang, and Shuaikang Hao. 2025. AI in the spotlight: The impact of artificial intelligence disclosure on user engagement in short-form videos. *Computers in Human Behavior* 162 (2025), 108448. doi:10.1016/j.chb.2024.108448

[3] Di Cooke, Abigail Edwards, Sophia Barkoff, and Kathryn Kelly. 2025. As Good As A Coin Toss: Human detection of AI-generated images, videos, audio, and audiovisual stimuli. arXiv:2403.16760 [cs.HC] https://arxiv.org/abs/2403.16760

[4] Abdallah El Ali, Karthikeya Puttur Venkatraj, Sophie Morosoli, Laurens Naudts, Natali Helberger, and Pablo Cesar. 2024. Transparent AI Disclosure Obligations: Who, What, When, Where, Why, How. In *Extended Abstracts of the CHI Conference on Human Factors in Computing Systems* (Honolulu, HI, USA) *(CHI EA '24)*. Association for Computing Machinery, New York, NY, USA, Article 342, 11 pages. doi:10.1145/3613905.3650750

[5] Ziv Epstein, Antonio Alonso Arechar, and David Rand. 2023. What label should be applied to content produced by generative AI? (2023).

[6] Dilrukshi Gamage, Dilki Sewwandi, Min Zhang, and Arosha K Bandara. 2025. Labeling Synthetic Content: User Perceptions of Label Designs for AI-Generated Content on Social Media. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems (CHI '25)*. ACM, 1–29. doi:10.1145/3706598.3713171

[7] Maurice Jakesch, Megan French, Xiao Ma, Jeffrey T. Hancock, and Mor Naaman. 2019. AI-Mediated Communication: How the Perception that Profile Text was Written by AI Affects Trustworthiness. In *Proceedings of the 2019 CHI Conference on Human Factors in Computing Systems* (Glasgow, Scotland Uk) *(CHI '19)*. Association for Computing Machinery, New York, NY, USA, 1–13. doi:10.1145/3290605.3300469

[8] Zhuoyan Li, Chen Liang, Jing Peng, and Ming Yin. 2024. How Does the Disclosure of AI Assistance Affect the Perceptions of Writing? arXiv:2410.04545 [cs.CL] https://arxiv.org/abs/2410.04545

[9] Sue Lim and Ralf Schmälzle. 2024. The effect of source disclosure on evaluation of AI-generated messages. *Computers in Human Behavior: Artificial Humans* 2, 1 (2024), 100058. doi:10.1016/j.chbah.2024.100058

[10] Chiara Longoni, Andrey Fradkin, Luca Cian, and Gordon Pennycook. 2022. News from Generative Artificial Intelligence Is Believed Less. In *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency* (Seoul, Republic of Korea) *(FAccT '22)*. Association for Computing Machinery, New York, NY, USA, 97–106. doi:10.1145/3531146.3533077

[11] Gordon Pennycook, Adam Bear, Evan T Collins, and David G Rand. 2020. The implied truth effect: Attaching warnings to a subset of fake news headlines increases perceived accuracy of headlines without warnings. *Management science* 66, 11 (2020), 4944–4957.

[12] Irene Rae. 2024. The Effects of Perceived AI Use On Content Perceptions. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems* (Honolulu, HI, USA) *(CHI '24)*. Association for Computing Machinery, New York, NY, USA, Article 978, 14 pages. doi:10.1145/3613904.3642076

[13] Manav Raj, Justin Berg, and Rob Seamans. 2024. Art-ificial Intelligence: The Effect of AI Disclosure on Evaluations of Creative Content. arXiv:2303.06217 [cs.CY] https://arxiv.org/abs/2303.06217

[14] David B. Resnik and Mohammad Hosseini and. 2025. Disclosing artificial intelligence use in scientific research and publication: When should disclosure be mandatory, optional, or unnecessary? *Accountability in Research* 0, 0 (2025), 1–13. doi:10.1080/08989621.2025.2481949 arXiv:https://doi.org/10.1080/08989621.2025.2481949 PMID: 40126451.

[15] E Union. 2021. Proposal for a regulation of the european parliament and of the council laying down harmonised rules on artificial intelligence (artificial intelligence act) and amending certain union legislative acts. *COM/2021/206final* (2021).

[16] U.S. Congress. 2023. AI Disclosure Act of 2023. https://www.congress.gov/bill/118th-congress/house-bill/3831. Accessed: 2025-04-24.

A Appendix

Content	Item	Responses	Percentage
Email	**Only indicate whether AI was used**	**18**	**54.5%**
	Specify AI's contribution numerically	1	3.0%
	Specify which part AI was used in	2	6.1%
	Specify which stage AI was used in	7	21.2%
	Other	6	18.2%
Personal Statement	**Only indicate whether AI was used**	**15**	**45.5%**
	Specify AI's contribution numerically	7	21.2%
	Specify which part AI was used in	8	24.2%
	Specify which stage AI was used in	13	39.4%
	Other	3	9.1%
Social Media Post	**Only indicate whether AI was used**	**15**	**45.5%**
	Specify AI's contribution numerically	6	18.2%
	Specify which part AI was used in	6	18.2%
	Specify which stage AI was used in	13	39.4%
	Other	3	9.1%

Table 3: Results of Survey on Preference for Simple AI Disclosure (Email, Social Media Post, Personal Statement)

Content	Item	Responses	Percentage
News	Only indicate whether AI was used	7	21.2%
	Specify AI's contribution numerically	6	18.2%
	Specify which part AI was used in	11	33.3%
	Specify which stage AI was used in	**23**	**69.7%**
	Other	1	3.0%
Textbook	Only indicate whether AI was used	8	24.2%
	Specify AI's contribution numerically	11	33.3%
	Specify which part AI was used in	16	48.5%
	Specify which stage AI was used in	**26**	**78.8%**
	Other	0	0.0%
Novel	Only indicate whether AI was used	7	21.2%
	Specify AI's contribution numerically	8	24.2%
	Specify which part AI was used in	12	36.4%
	Specify which stage AI was used in	**23**	**69.7%**
	Other	2	6.1%
Assignment (Essay)	Only indicate whether AI was used	9	27.3%
	Specify AI's contribution numerically	10	30.3%
	Specify which part AI was used in	11	33.3%
	Specify which stage AI was used in	**26**	**78.8%**
	Other	0	0.0%

Table 4: Results of Survey on Preference for Detailed AI Disclosure (News, Novel, Textbook, Assignment(Essay))

		Necessary	Unnecessary
Creator	Want to disclose	5	1
	Don't want to disclose	3	25
Reader	Want disclosure	9	7
	Don't want disclosure	1	17

(a) Email

		Necessary	Unnecessary
Creator	Want to disclose	25	1
	Don't want to disclose	7	1
Reader	Want disclosure	31	1
	Don't want disclosure	1	1

(b) Textbook

Table 5: Strong match contents, where creators' and readers' perspectives on AI disclosure obligations and preferences align

		Necessary	Unnecessary
Creator	Want to disclose	16	1
	Don't want to disclose	14	4
Reader	Want disclosure	28	2
	Don't want disclosure	2	2

(a) News

		Necessary	Unnecessary
Creator	Want to disclose	11	2
	Don't want to disclose	15	5
Reader	Want disclosure	26	4
	Don't want disclosure	1	3

(b) Assignment (Essay)

Table 6: Medium match contents, where creators' and readers' perspectives on AI disclosure obligations and preferences partially align

		Necessary	Unnecessary
Creator	Want to disclose	12	4
	Don't want to disclose	6	12
Reader	Want disclosure	22	5
	Don't want disclosure	1	6

(a) Social Media Post

		Necessary	Unnecessary
Creator	Want to disclose	6	3
	Don't want to disclose	10	14
Reader	Want disclosure	21	5
	Don't want disclosure	1	7

(b) Personal Statement

		Necessary	Unnecessary
Creator	Want to disclose	13	3
	Don't want to disclose	7	10
Reader	Want disclosure	20	5
	Don't want disclosure	5	4

(c) Novel

Table 7: Low match contents, where creators' and readers' perspectives on AI disclosure obligations and preferences diverge